

# CQRS / EventSourcing

Softwerkskammer Leipzig

2016-10-18

# Agenda

- CQRS in 10 mins
- EventSourcing in 10 mins
- Q&A
- Workshop intro
- Workshop (yes, **you** code) 1-3
- break
- Workshop (yes, **you** code) 4-7
- Q&A



... and yes, WE HIRE!

# *Kudos*

Oliver Wolf

<https://innoq.com>



Greg Young

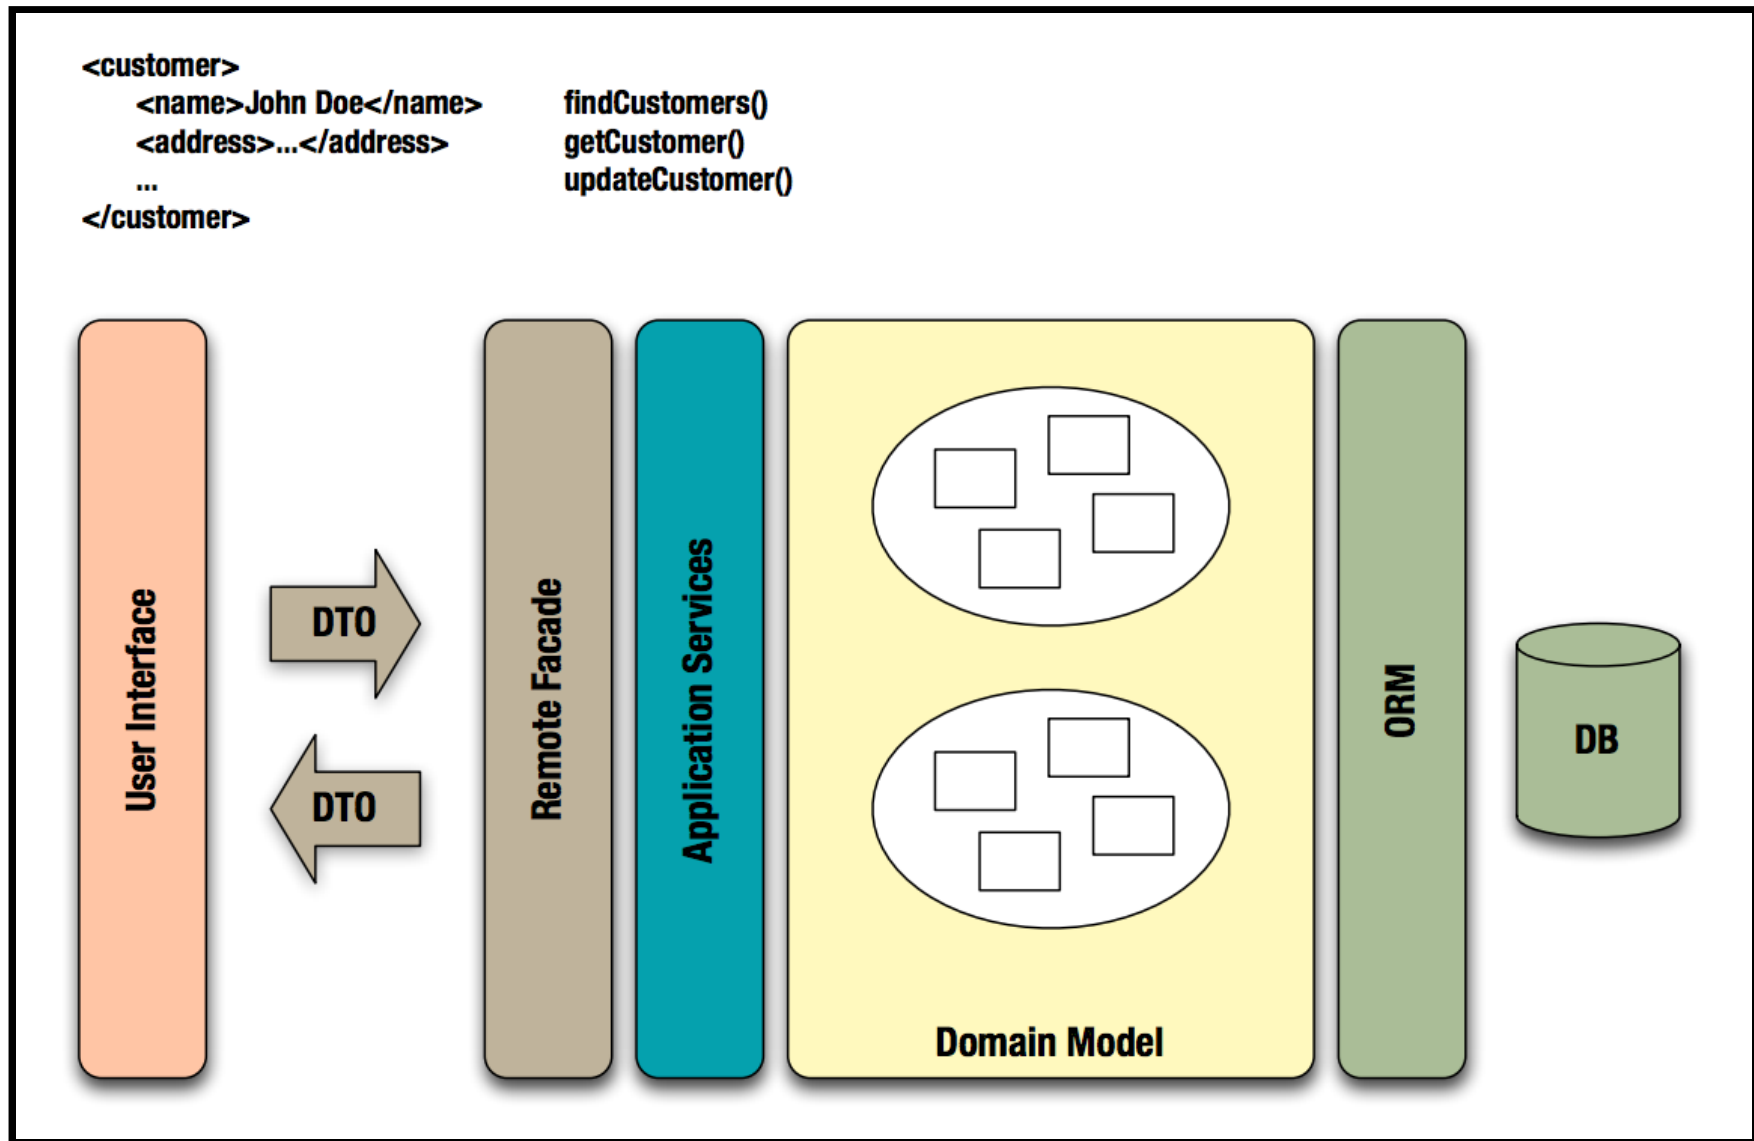
<https://goodenoughsoftware.net/>



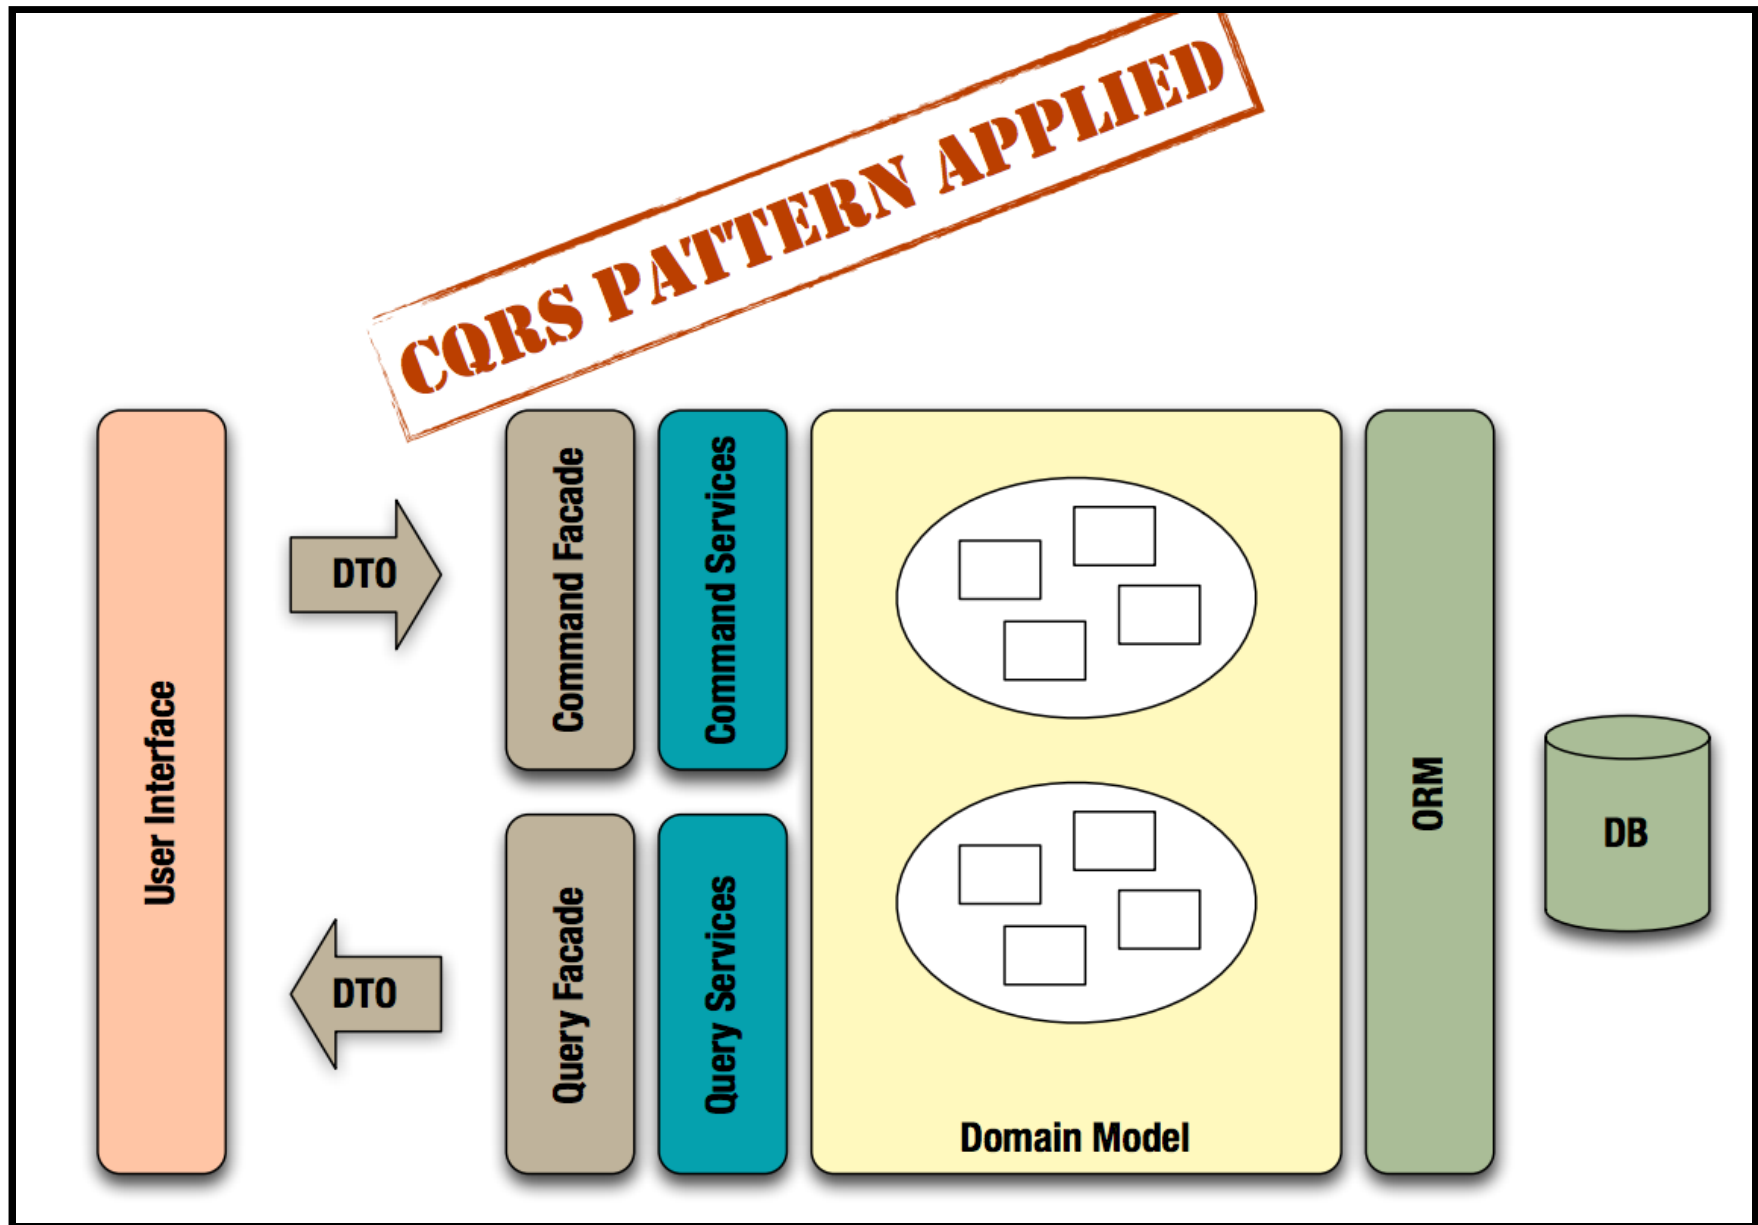
"You can do CQRS without EventSourcing, but you cannot do EventSourcing without CQRS."

# CQRS

COMMAND QUERY RESPONSIBILITY SEGREGATION



Traditional Layered Architecture

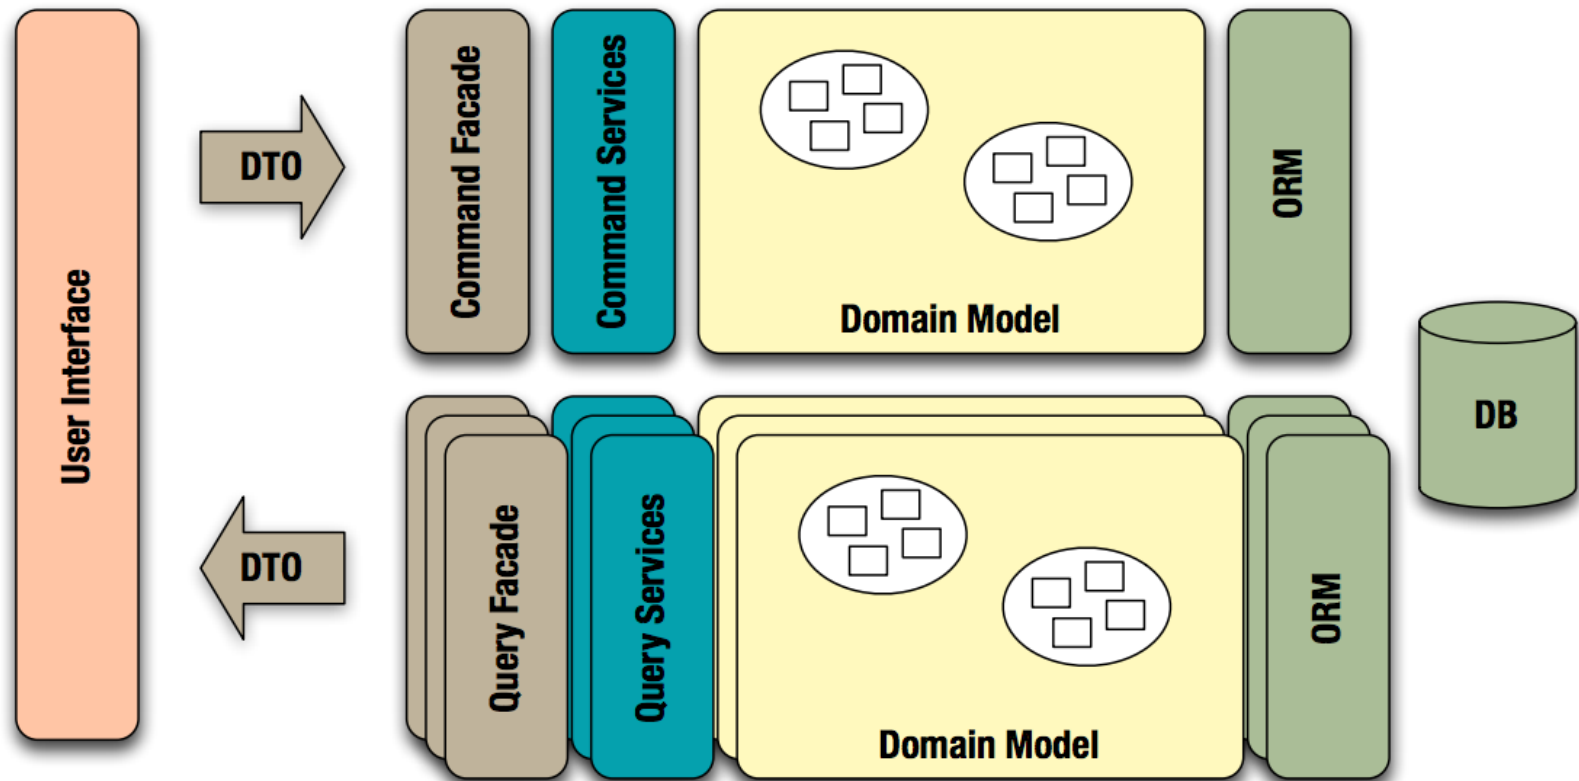


CQRS Pattern applied – done.



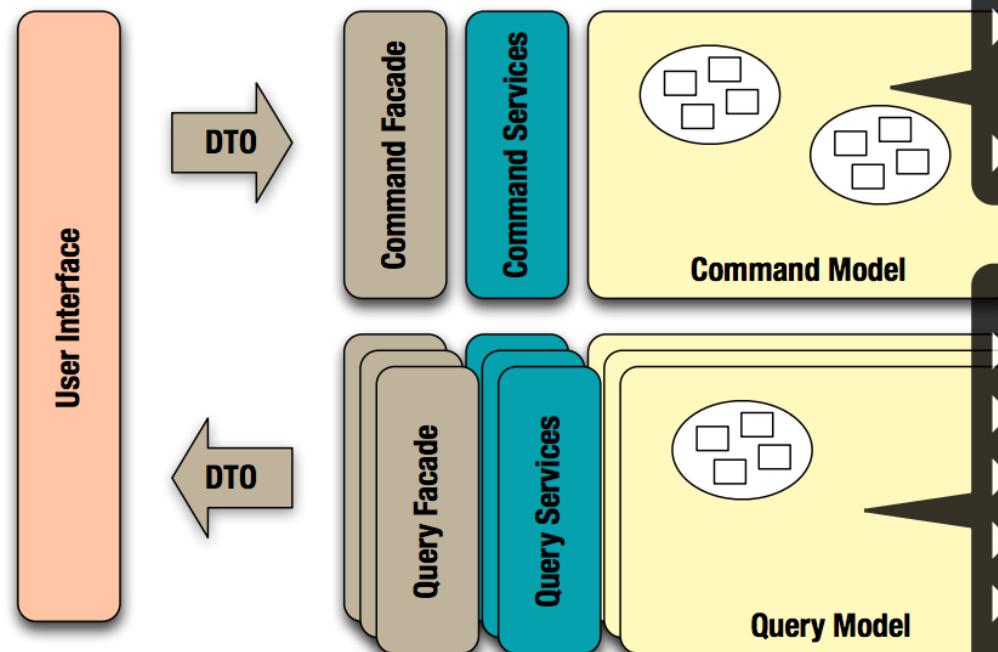
THAT IS IT!?

**Command and query parts can scale independently, e.g. to accommodate highly asymmetric load.**



pro: scale independently

**Queries can benefit from a specialized query model, optimized for quick data retrieval (de-normalized, pre-aggregated, ...)**




- ▶ validate and process commands
- ▶ keep data consistent
- ▶ guarantee ACID properties
- ▶ behaviour part of domain model
- ▶ relatively difficult to scale out

- ▶ rich query capabilities
- ▶ short response times
- ▶ different views on data
- ▶ potentially denormalized
- ▶ relatively easy to scale out

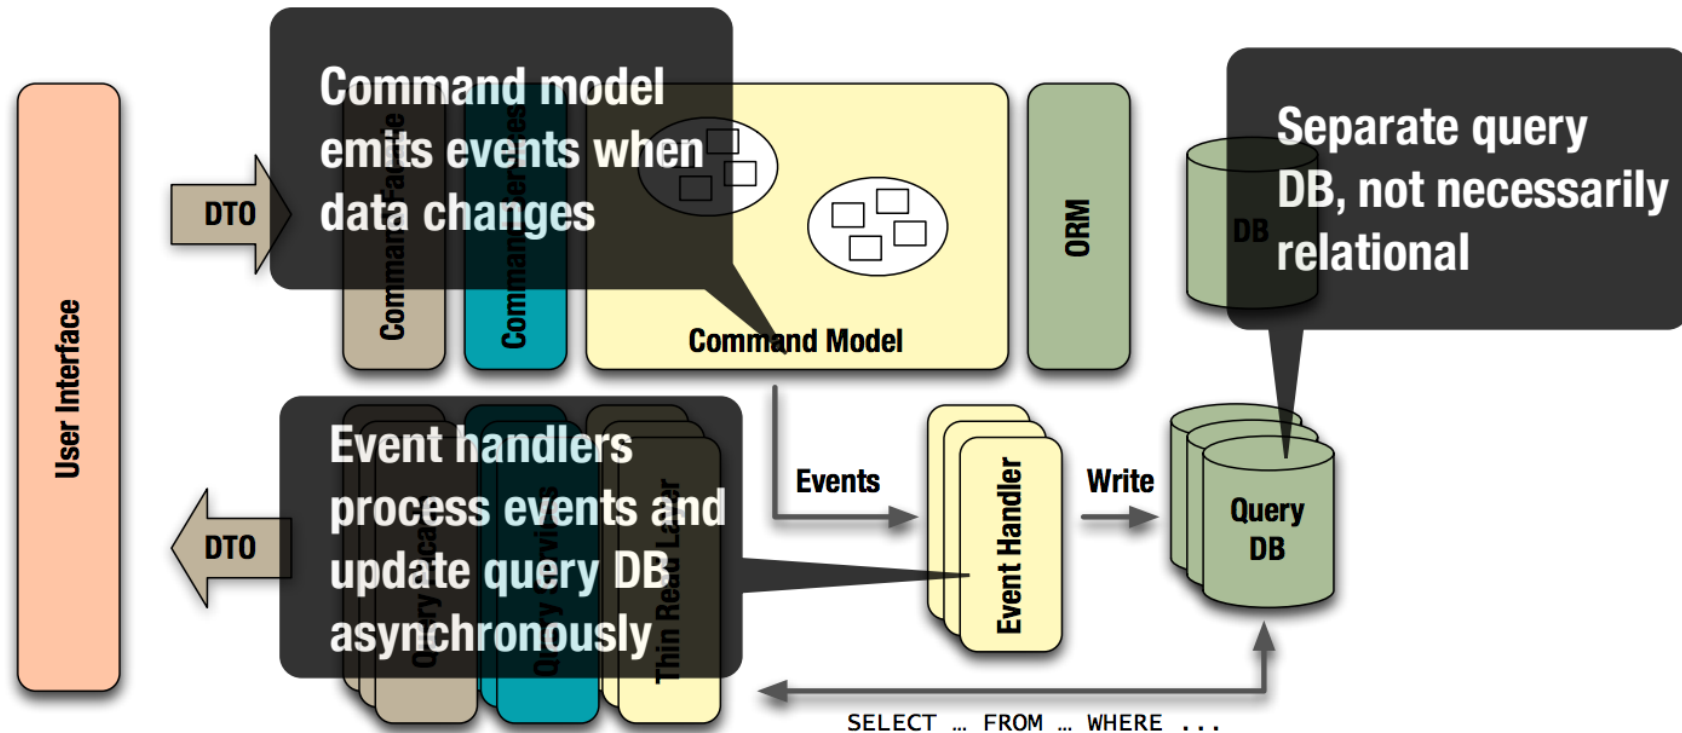
pro: optimized Query-Models (denormalization)

Queries are just dealing with data, not with behaviour.  
What do we need objects for?



pro: **thin** read layer

In many cases, **eventual consistency** is sufficient.  
The data users are looking at in the UI is always stale to some extent.



pro: Eventual Consistency & Read-Replicas

# Conclusion

CQRS helps with

- asymmetric load / read replicas
- gain from different QueryModels / Technologies
- helps Time-to-Market
- avoids technology Lock-In
- enables local optimization on Query-Models

# EventSourcing

as a Concept

"An architectural pattern which warrants that your entities (as per Eric Evans' definition) do not track their internal state by means of *direct serialization or O/R mapping*, but by means of **reading and committing events** to an event store."

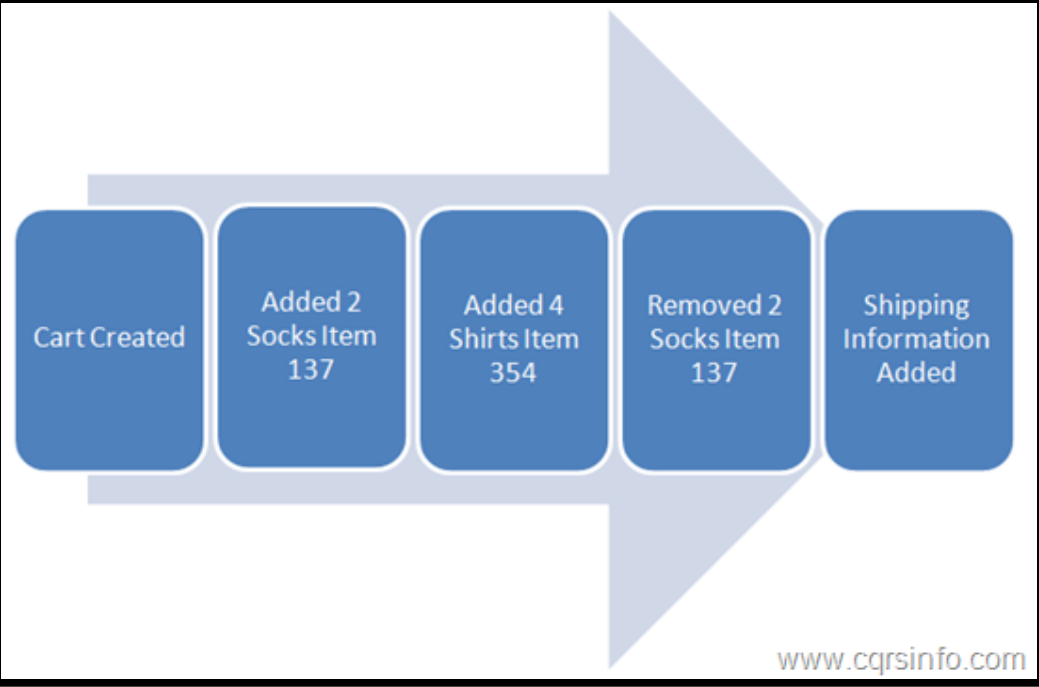


# EventSourcing

- Architectural Pattern
- EventStore keeps log of Events (Facts)
- 'current' State is
  - transient
  - disposable
  - fully/reliably reconstructible from series of Events

## Current State is a Left Fold of Events

- FP: Left Fold aggregates a collection via a function and an initial value
  - Ruby: `[1, 2, 3].inject(0, :+)` == 6 # symbol fn name
  - Scala: `List(1, 2, 3).foldLeft(0)(_ + _)` == 6 // anon function
- Provide an initial state  $s_0$  and a function  $f : (S, E) \Rightarrow S$
- Current State after event  $e_3$  is:
  - = `leftFold( [e1, e2, e3], s0, f )`
  - = `f( f( f( s0, e1), e2), e3 )`



# Pros

- focus on **state transitions**, rather than data structure
- audit log already included
- reports over the past
- no Information-loss (Socks, Item 137)
- replayable
  - basically Time Machine (travel back and forth)
  - history of System state
  - helps debugging
  - no infamous SQL-migration-scripts, just change aggregation and replay

# Cons

- maybe different angle to modeling
- (little) more complex than a CRUD System
- new Challenges like:
  - aggregation performance
  - evolving events
  - capacity

proven Patterns for the new challenges **do** exist!

# Myths

- artificial approach to modeling (**not** true)
- requires eventual consistent (**not** true)
- inherently difficult & complex (**not** true)
- bad performance (**not** true)

Questions up to here?

EventSourcing Basics

# Workshop



# Disclaimer

the code here is

- **NOT** an EventSourcing framework !
- just for demonstration of concepts
- oversimplified
- Java, but can be done in any language
- very basic DI with Spring, but can be done without
- uses Lombok for brevity (just syntactical sugar)



There are many ways to skin a Cat

This Workshop is about  
**Discussion**  
not Code

EventSourcing Basics


# Intro

Hello, we are FooBank !

# Our Domain

- Local Bank
- physical Counter
- will expand into online-banking

# Infrastructure



# Component      Responsibility

## Glossary

Component	Responsibility
ApplicationFacade	single Entry Point / internal API (optional)
Command	Request for the System to <i>do</i> something
CommandBus	find CmdHandler for given Cmd
CommandHandler	accept or reject Command, emit Effects
Effects	List of Messages
Message	Event or Command
Event	a given Fact
EventStore	a log of Events <i>that have happened</i>
EventHandler/View	process Events, project useful Model
Query	a Question to a Model
QueryModel	queried by the outside world, query-optimized data
ValidationModel	answers Queries while validating



EventSourcing Basics

Session 1

# Aggregating to the canonical Domain Model

# Canonical Domain Model

```
public class Account {  
    private final UUID id;  
    private final String firstName;  
    private final String lastName;  
    private int balance = 0;  
  
    void credit(int amount) { balance += amount; }  
    void debit(int amount) { balance -= amount; }  
}
```

# UseCase Deposition

## UseCase Deposition

As a Customer


i want to **deposit** cash at the counter  
in order to credit it to my account.

# UseCase Withdrawal

## UseCase Withdrawal

As a Customer

i want to **withdraw** money from my account at the counter  
in order to cash it out.



# What we learn

- Implement **write** side
  - minimal Commands
  - minimal CommandHandlers
  - minimal Events
- Implement **read** side
  - minimal EventHandlers (Views)
  - that populate the canonical Domain Model

```
git clone https://github.com/uweschaefer/es-basics.git
```

# Session 1

1. Implement **ApplicationFacade.deposit/withdraw**
2. Create **Command classes** for both UseCases (see *CreateAccountCommand*)
3. Create **CommandHandlers** for both UseCases (see *CreateAccountHandler*)
4. Create **Event** classes for every UseCase (see *AccountCreatedEvent*)
5. Extend **AccountView** to aggregate Accounts
6. Pass the Tests



# What just happened?

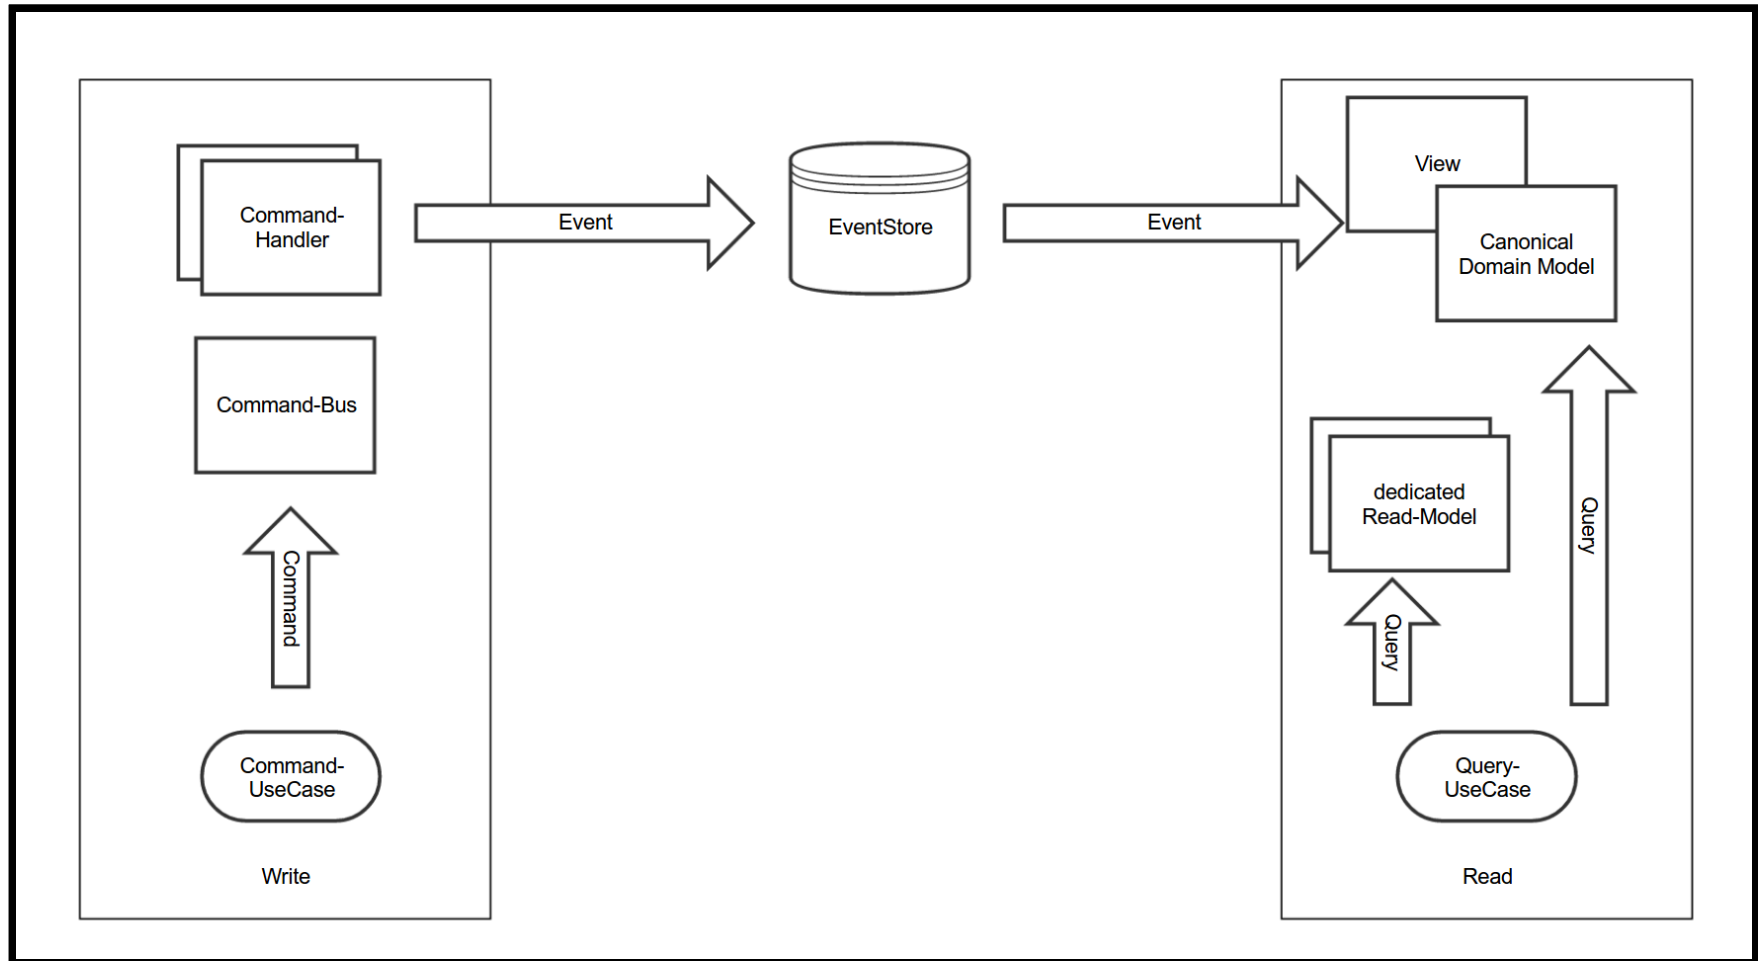
- complete Roundtrip
  - Command -> CommandHandler
  - Event -> EventHandler (View)
  - Query via DomainModel

EventSourcing Basics

Session 2

# Dedicated QueryModel

# Overview



# Aggregation

AccountView = EventHandler that aggregates Events to an Account.

Not every Query within one Account

## UseCase ValuedCustomerReport

## UseCase ValuedCustomerReport

As a Manager, i want  
a complete report the lists all *valued customers*  
in order to free them from handling charges.

**Specification ValuedCustomer**

**Specification ValuedCustomer**

*a valued Customer*

has deposited an amount of  **$\geq 1000\text{€}$  at least twice.**

# How **not** to do that

Iterate Accounts and inspect their Depositions one by one.



# What we learn

- Aggregate beyond Entity-Boundaries
- Create **dedicated Read Model**
  - find appropriate DataStructure
  - select Event-Types by @EventConsumer Methods

# Session 2

1. `git clean -fd && git reset --hard`  
session2
2. implement  
ValuedCustomerReportView
3. pass the Tests

# What just happened?

- we added a dedicated Read/Query-Model
  - beyond aggregate boundaries
  - Query-optimized Datastructure
  - PullViews have to actually **pull** the events from the ES at some point.

EventSourcing Basics  
Session 3 (Bonuslevel)

# Rolling Snapshot QueryModel

# Overview

- ValuedCustomerSupport aggregates ALL depositions in the System
- gets slower & slower
- how to tackle that?

# Problematic

every time a Report is needed, a new View has to be created.

```
facade.deposit(...);  
  
ValuedCustomerReportView report1 = new ValuedCustomerReportView(es);  
assertTrue(report1.isValuedCustomer(...));  
  
facade.deposit(...);  
  
ValuedCustomerReportView report2 = new ValuedCustomerReportView(es);  
assertTrue(report2.isValuedCustomer(...));
```

# What if?

we could reuse a QueryModel, that is being updated, rather than re-created?

```
ValuedCustomerReportView report = new ValuedCustomerReportView(es);  
  
facade.deposit(...);  
assertTrue(report.isValuedCustomer(...));  
  
facade.deposit(...);  
assertTrue(report.isValuedCustomer(...));
```

---

**discuss** what would be necessary, conceptually?

No, really – **Discuss**



# What we just learned?

- Concept of Rolling Snapshot

# Session 3

1. `git clean -fd && git reset --hard session3`
2. look at *View.last*, *View.accept* and *PullView.pullEvents*
3. change `ValueCustomerReport` appropriately.
4. pass the Tests

# What just happened?

- Rolling Snapshot
  - keeps latest
- Eventstore provides query of EventStream from **after** a particular event
- And yes, we can have more snapshots than one, if needed

Have a break.



Have a KitKat

EventSourcing Basics

Session 4

# Event Design

# UseCase Transfer

## UseCase Transfer

As a user, i want to  
transfer Money from my account to someone else's  
in order to pay my rent online.

# Acceptance Criteria

- *AccountUnknownException* if receiver or sender account does not exist
- *UnfundedTransferException* if sender does not have enough money (no debt allowed)

# What we learn

- Event granularity matters
- Events need to reveal their intent
- Use of a ValidationModel
- Commands can be rejected



# Session 4

1. git clean -fd && git reset --hard  
session4
2. implement *TransferHandler*
3. apply necessary changes AccountView
4. pass the Tests

# What just happened?

- Granularity: Events belong to ONE Aggregate
  - we need SendTransfer, RecieveTransfer
- Events reveal intent
  - do not reuse WithdrawnEvent etc, its a different UseCase!

EventSourcing Basics

Session 5 (Bonuslevel)

# Dedicated WriteModel / ValidationModel

Command validation sometimes needs Context

# Your take on Criteria 1?

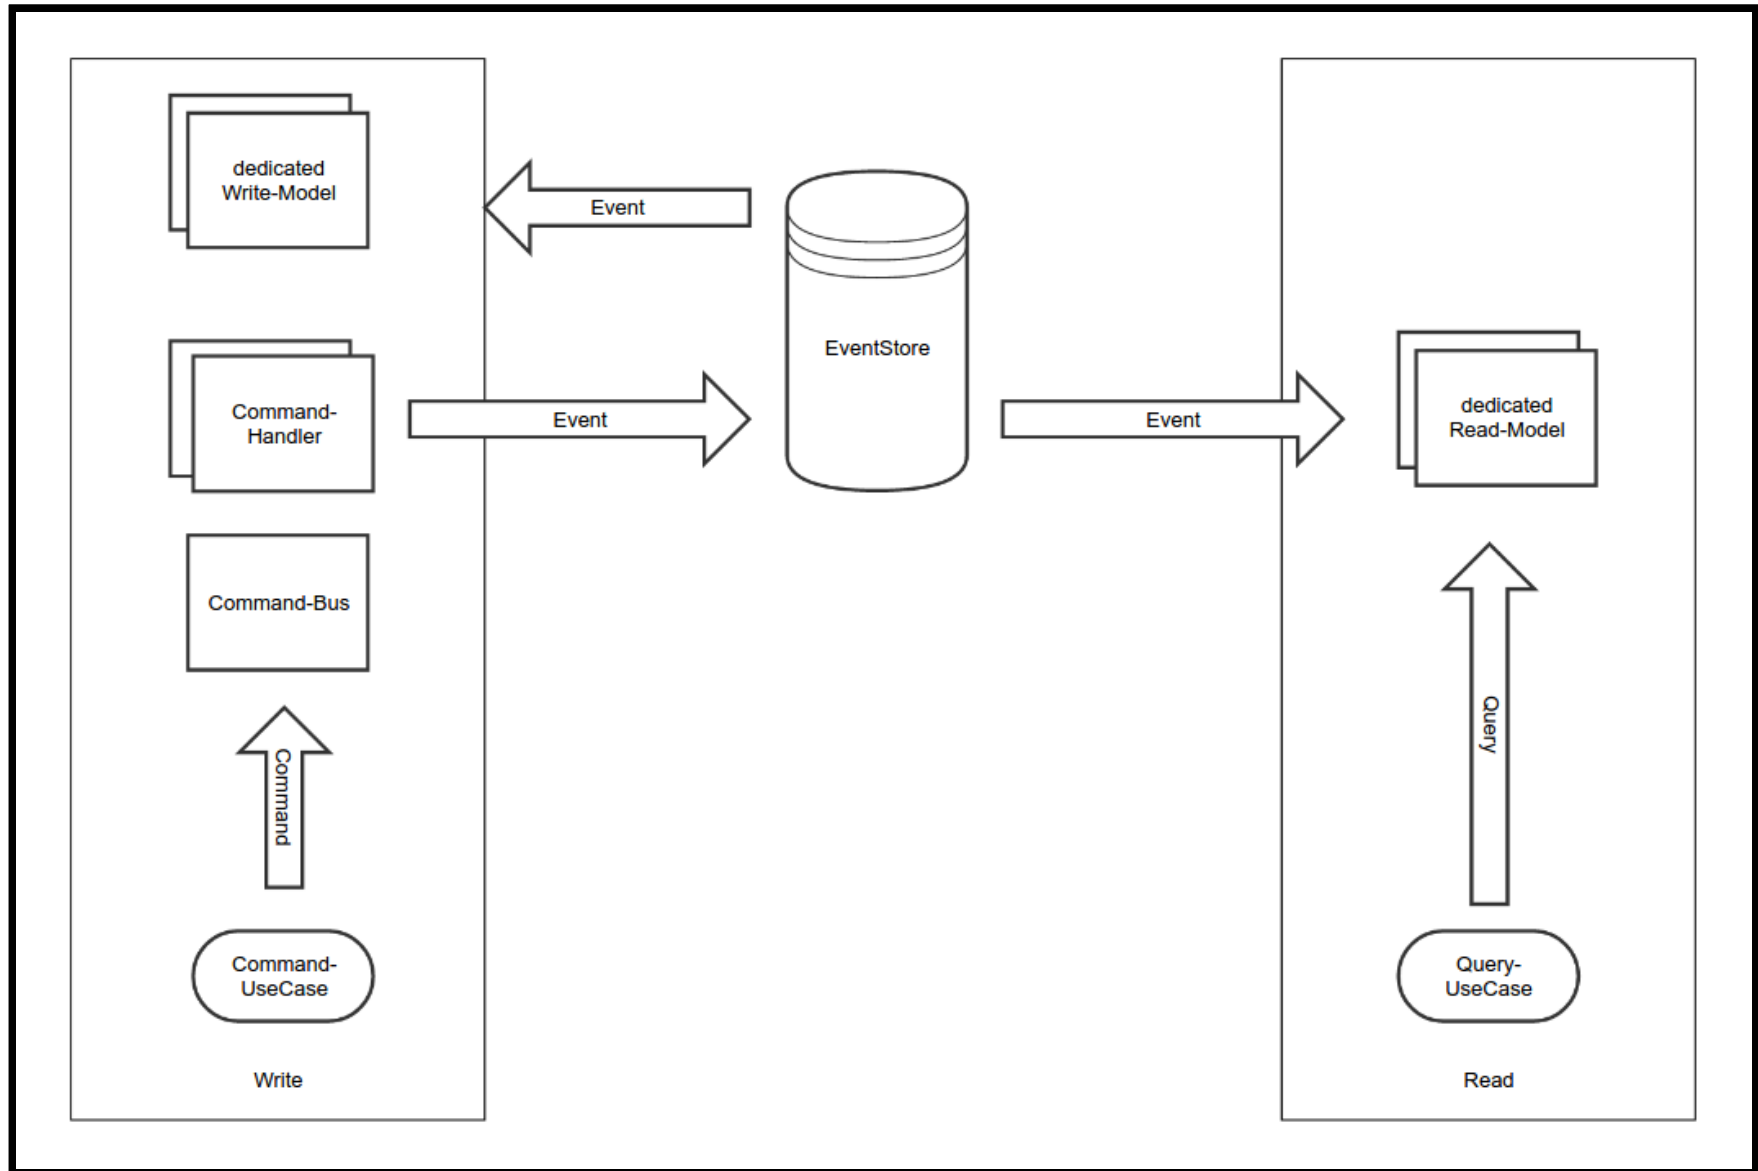
- *AccountUnknownException* if receiver or sender account does not exist

```
private boolean exists(UUID id) {  
    return repo.find(id) != null;  
}
```

Using AccountView just to find out, if an Account exists is wasteful.

All we really need to know  
is if the aggregate exists.

# Overview



# CommandHandler

## Responsibilities

- validate command
- accept or reject command based on that validation
- emit Messages on accepting



# Session 5

1. `git clean -fd && git reset --hard session5`
2. implement `KnownAccountsView`
3. pass the Tests

# What just happened?

- dedicated WriteModel / ValidationModel
- does not have to be DomainModel, as it does not need behaviour

EventSourcing Basics

Session 6

# Side-Effects

- Some Commands may trigger external behavior.
- Replaying that would be problematic.

# UseCase Notification

## UseCase Notification

As a user, i want to  
be notified by email when i recieve a transfer  
in order to buy champagne asap.

# What we learn

- how/where to model Side-Effect

# Session 6

1. `git clean -fd && git reset --hard session6`
2. use *CreditNotificationService* to send mail
3. discuss where/how to do it properly
4. hint: see *CommandBus.publish()*
5. pass the Tests

# What just happened?

- CommandBus has to be reliable
- Commands can be Effects, too
- Side-Effects can be modeled as Commands / CommandHandlers



# Intermission

What about Consistency?

Did we relax Consistency compared to a normal CRUD/ORM implementation?

**NO!**

But where we could, how can we take advantage?

EventSourcing Basics  
Session 7 (Bonuslevel)

# Push-Views

Up to now, all views have been *PullViews*, that call *pullEvents()* to stream events into them.

# Pro

- we can define when to update the View's State



# Con

- we have to Query the EventStore in order to know, if View's State is stale
  - the more Queries we run, the more catastrophic this is:
  - bad Latency for Queries
  - high Contention on EventStore

## **UseCase GoldCustomers**

## **UseCase GoldCustomers**

As an accountant, i want to  
know all the Gold-Customers  
in order to be extra nice to them.

# Specification Gold Customer

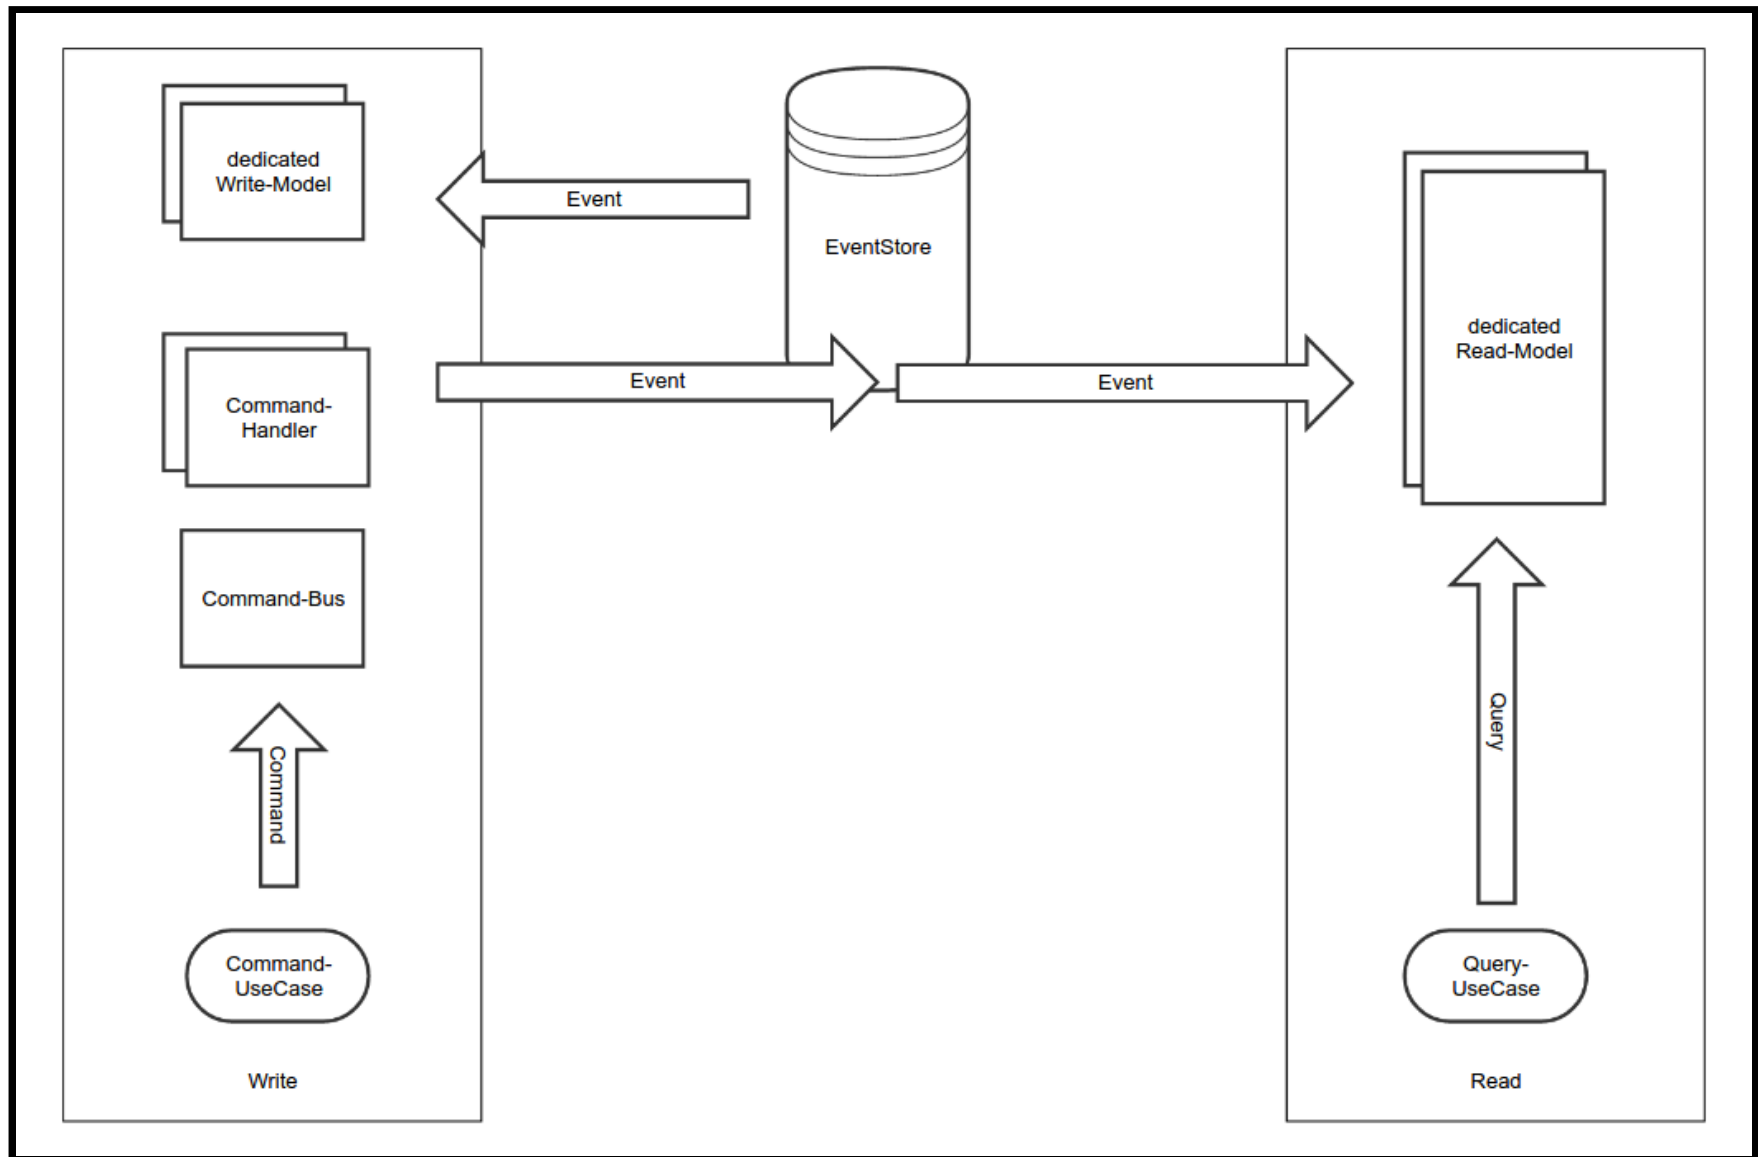
## Specification Gold Customer

someone who recieved a transfer  $\geq 10.000\text{€}$   
at least once

# Acceptance Criteria

- report must be **instant!** (low-latency)
- report must be a collection of Strings "<LASTNAME>, <FIRSTNAME>"
- order is not important
- only Transfers count – Depositions **must not** be examined
- report does not need to include GoldCustomer that recieved the status in the last few seconds...

... which means **Eventual Consistency** is ok



What would be necessary to push events to the view?

# What we learn

- Use Push-Model for Views
- pros and cons of push vs pull



# Session 7

1. git clean -fd && git reset --hard session7
2. implement *GoldCustomersView* extends *PushView*
3. pass the Tests

# What just happened?

- Implemented a push-View that is updated by processing Events asynchronously
- Push reduces read latency
- introduces eventual consistency
- introduces concurrency
- PushViews mostly unusable as Validation Model (not strictly consistent)

One possible solution can be found here  
`git clean -fd && git reset --hard theend`

# Links

- O.Wolfs CQRS Slides
  - <https://speakerdeck.com/owolf/cqrs-for-great-good-2>
- Greg Young's Blog
  - <https://goodenoughsoftware.net/>
- Axon mature ES Framework
  - <http://www.axonframework.org/>
- Lagom Modern ES Framework based on Akka
  - <https://www.lightbend.com/lagom>
- Microsofts CQRS/ES Patterns & Practices
  - <https://msdn.microsoft.com/en-us/library/jj554200.aspx>



... and yes, WE HIRE!

Q & A