

CQRS / EventSourcing

Softwerkskammer Leipzig

2016-10-18

Agenda

- CQRS in 10 mins
- EventSourcing in 10 mins
- Q&A
- Workshop intro
- Workshop (yes, **you** code) 1-3
- break
- Workshop (yes, **you** code) 4-7
- Q&A

Kudos

Oliver Wolf

<https://innoq.com>



Greg Young


<https://goodenoughsoftware.net/>




"You can do CQRS without EventSourcing, but you cannot do EventSourcing without CQRS."

CQRS

COMMAND QUERY RESPONSIBILITY SEGREGATION




Traditional Layered Architecture



CQRS Pattern applied – done.

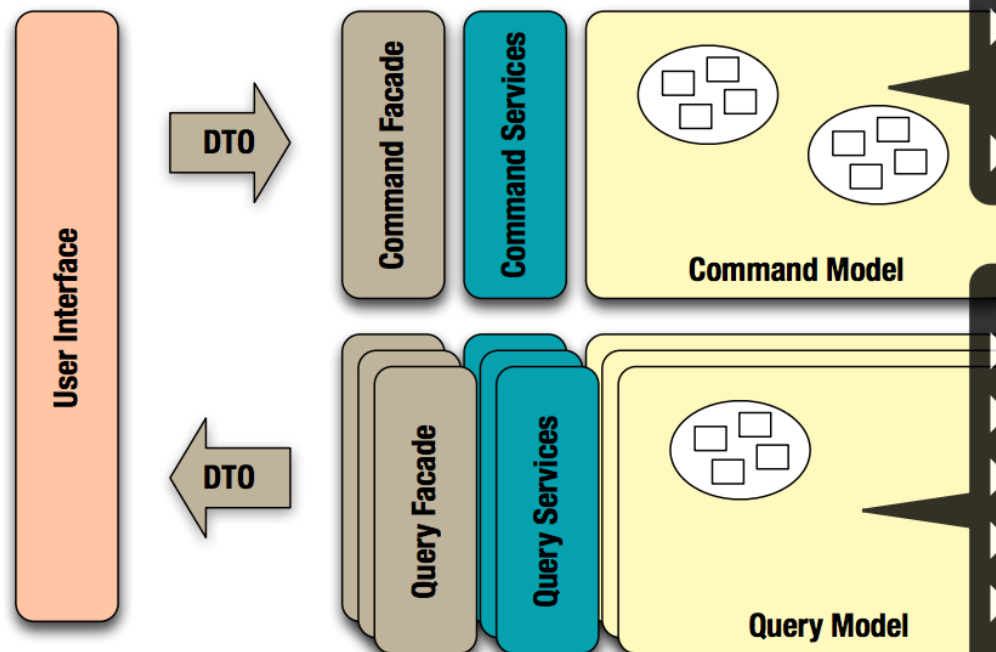
THAT IS IT!?

Command and query parts can scale independently, e.g. to accommodate highly asymmetric load.



pro: scale independently

Queries can benefit from a specialized query model, optimized for quick data retrieval (de-normalized, pre-aggregated, ...)

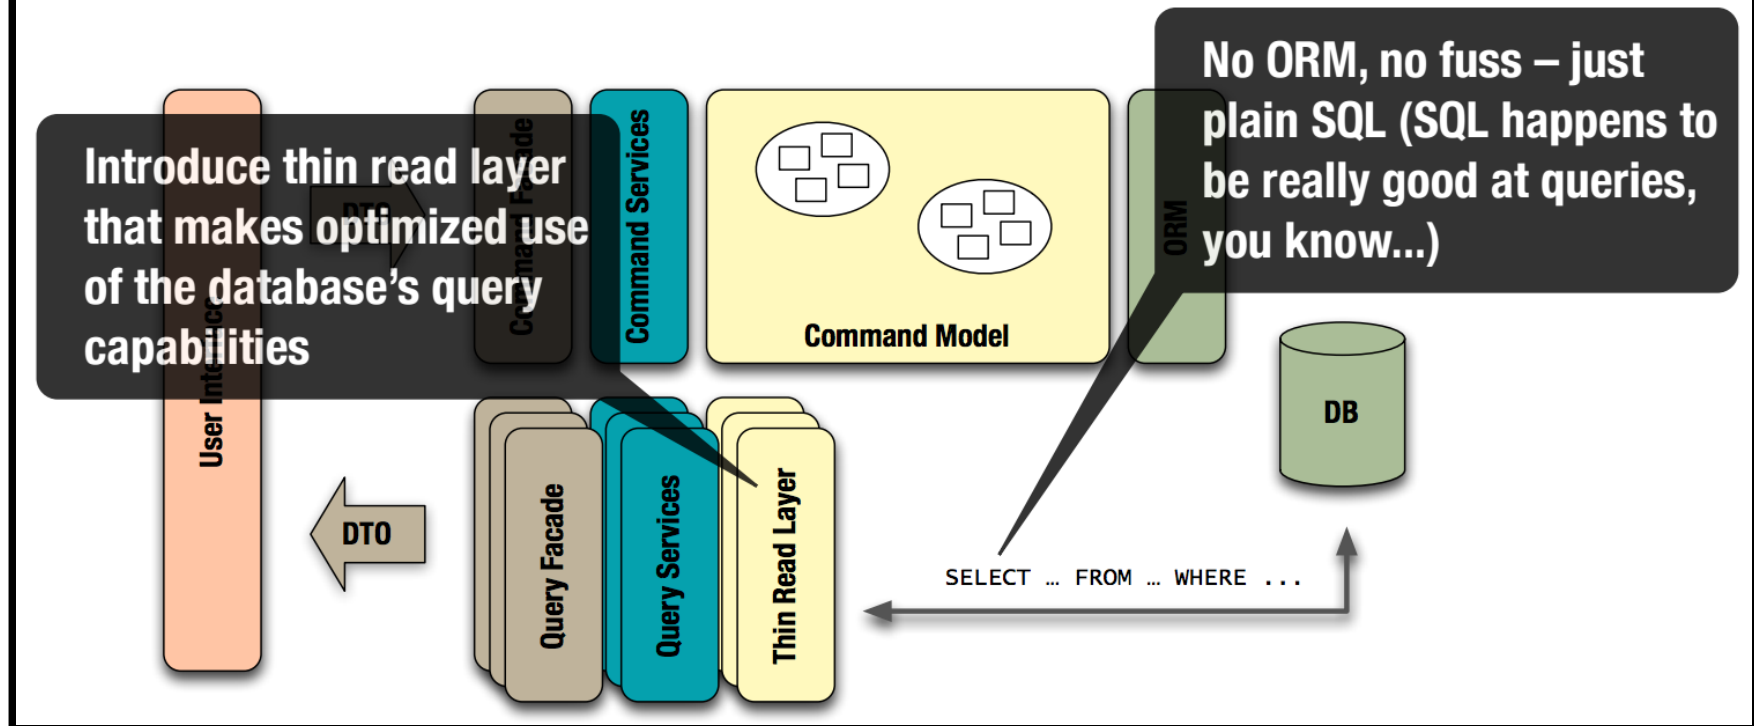


- ▶ validate and process commands
- ▶ keep data consistent
- ▶ guarantee ACID properties
- ▶ behaviour part of domain model
- ▶ relatively difficult to scale out

- ▶ rich query capabilities
- ▶ short response times
- ▶ different views on data
- ▶ potentially denormalized
- ▶ relatively easy to scale out


pro: optimized Query-Models (denormalization)

Queries are just dealing with data, not with behaviour.
What do we need objects for?



pro: **thin** read layer

In many cases, **eventual consistency** is sufficient.
The data users are looking at in the UI is always stale to some extent.



pro: Eventual Consistency & Read-Replicas

Conclusion

CQRS helps with

- asymmetric load / read replicas
- gain from different QueryModels / Technologies
- helps Time-to-Market
- avoids technology Lock-In
- enables local optimization on Query-Models

EventSourcing

as a Concept

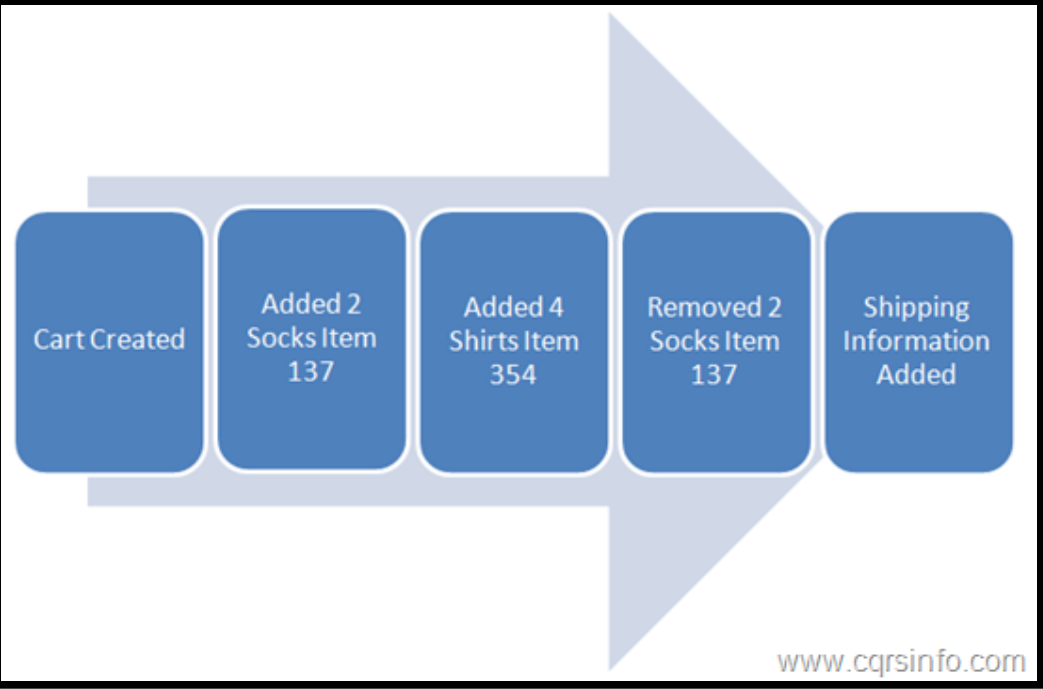
"An architectural pattern which warrants that your entities (as per Eric Evans' definition) do not track their internal state by means of *direct serialization or O/R mapping*, but by means of **reading and committing events** to an event store."

EventSourcing

- Architectural Pattern
- EventStore keeps log of Events (Facts)
- 'current' State is
 - transient
 - disposable
 - fully/reliably reconstructible from series of Events

Current State is a Left Fold of Events

- FP: Left Fold aggregates a collection via a function and an initial value
 - Ruby: `[1, 2, 3].inject(0, :+)` == 6 # symbol fn name
 - Scala: `List(1, 2, 3).foldLeft(0)(_ + _)` == 6 // anon function
- Provide an initial state s_0 and a function $f : (S, E) \Rightarrow S$
- Current State after event e_3 is:
 - = `leftFold([e1, e2, e3], s0, f)`
 - = `f(f(f(s0, e1), e2), e3)`



Pros

- focus on **state transitions**, rather than data structure
- audit log already included
- reports over the past
- no Information-loss (Socks, Item 137)
- replayable
 - basically Time Machine (travel back and forth)
 - history of System state
 - helps debugging
 - no infamous SQL-migration-scripts, just change aggregation and replay

Cons

- maybe different angle to modeling
- (little) more complex than a CRUD System
- new Challenges like:
 - aggregation performance
 - evolving events
 - capacity

proven Patterns for the new challenges **do** exist!

Myths

- artificial approach to modeling (**not** true)
- requires eventual consistent (**not** true)
- inherently difficult & complex (**not** true)
- bad performance (**not** true)

Questions up to here?

EventSourcing Basics

Workshop

Disclaimer

the code here is

- **NOT** an EventSourcing framework !
- just for demonstration of concepts
- oversimplified
- Java, but can be done in any language
- very basic DI with Spring, but can be done without
- uses Lombok for brevity (just syntactical sugar)



There are many ways to skin a Cat

This Workshop is about
Discussion
not Code

EventSourcing Basics

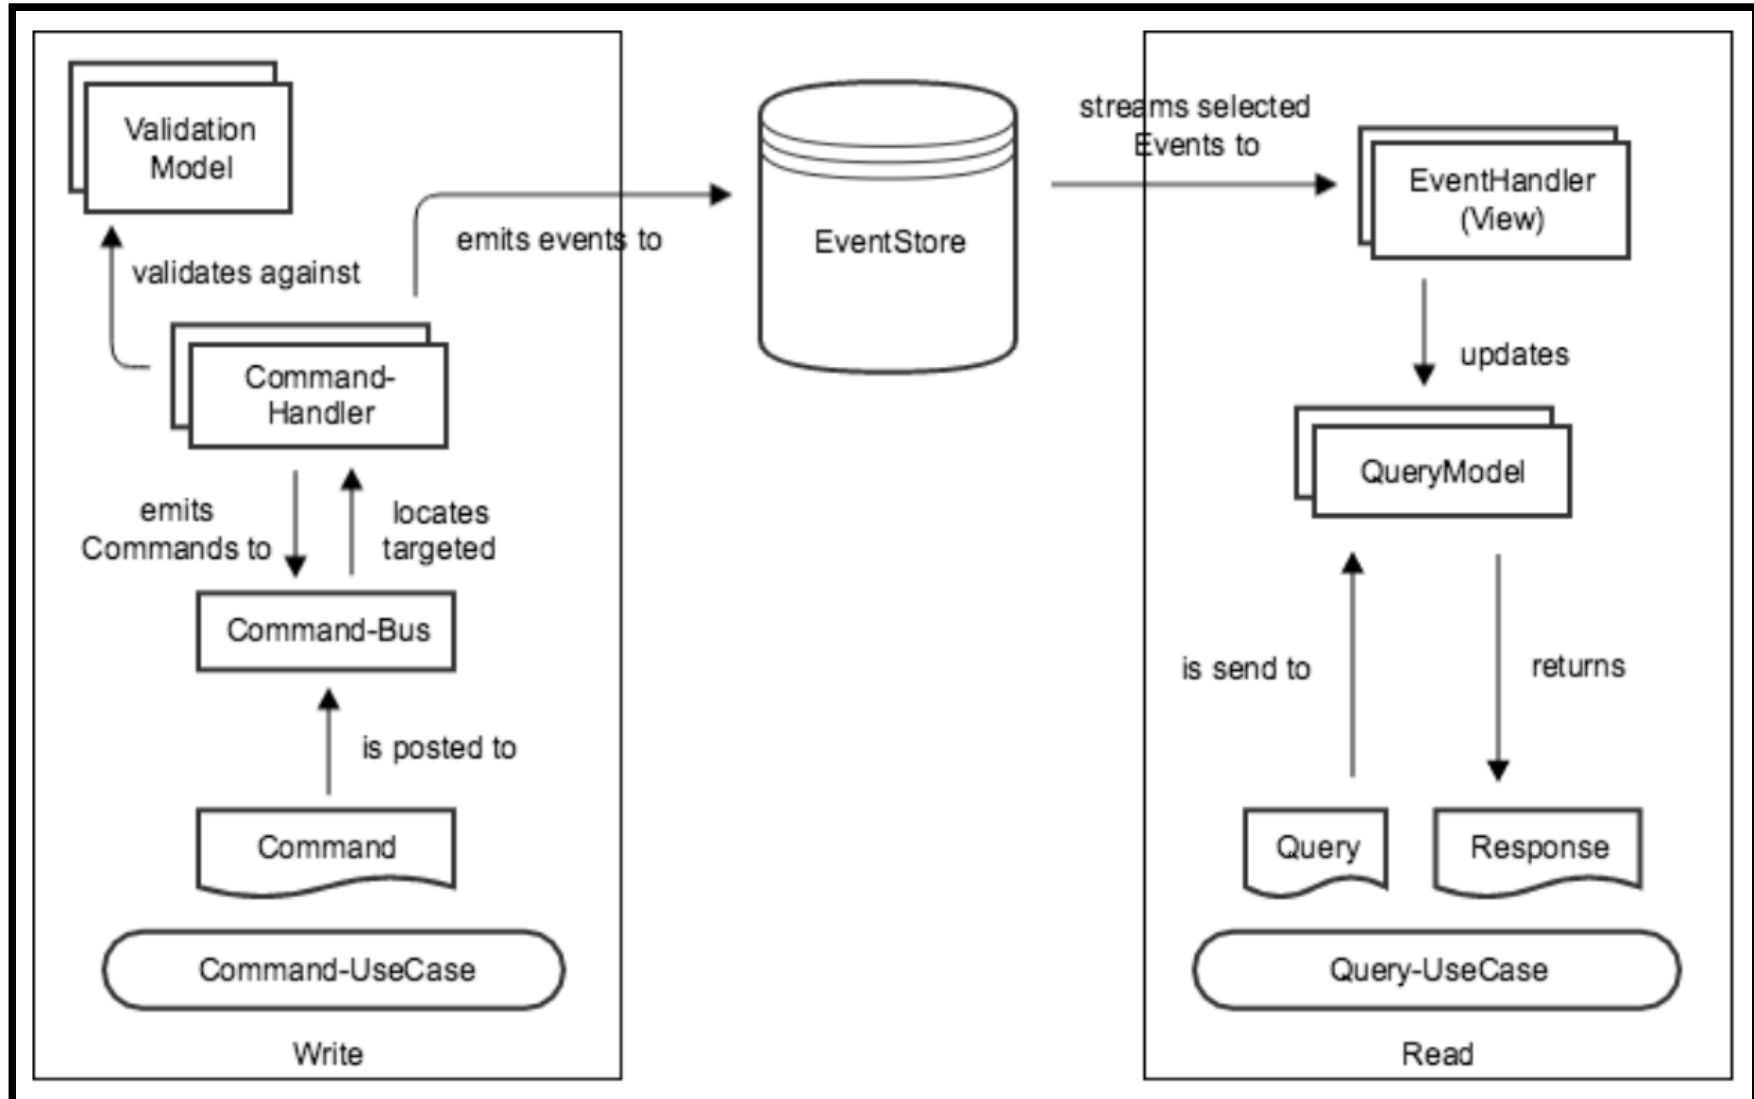
Intro

Hello, we are FooBank !

Our Domain

- Local Bank
- physical Counter
- will expand into online-banking

Infrastructure



Component Responsibility

Glossary

Component	Responsibility
ApplicationFacade	single Entry Point / internal API (optional)
Command	Request for the System to <i>do</i> something
CommandBus	find CmdHandler for given Cmd
CommandHandler	accept or reject Command, emit Effects
Effects	List of Messages
Message	Event or Command
Event	a given Fact
EventStore	a log of Events <i>that have happened</i>
EventHandler/View	process Events, project useful Model
Query	a Question to a Model
QueryModel	queried by the outside world, query-optimized data
ValidationModel	answers Queries while validating

EventSourcing Basics

Session 1

Aggregating to the canonical Domain Model

Canonical Domain Model

```
public class Account {  
    private final UUID id;  
    private final String firstName;  
    private final String lastName;  
    private int balance = 0;  
  
    void credit(int amount) { balance += amount; }  
    void debit(int amount) { balance -= amount; }  
}
```

UseCase Deposition

UseCase Deposition

As a Customer

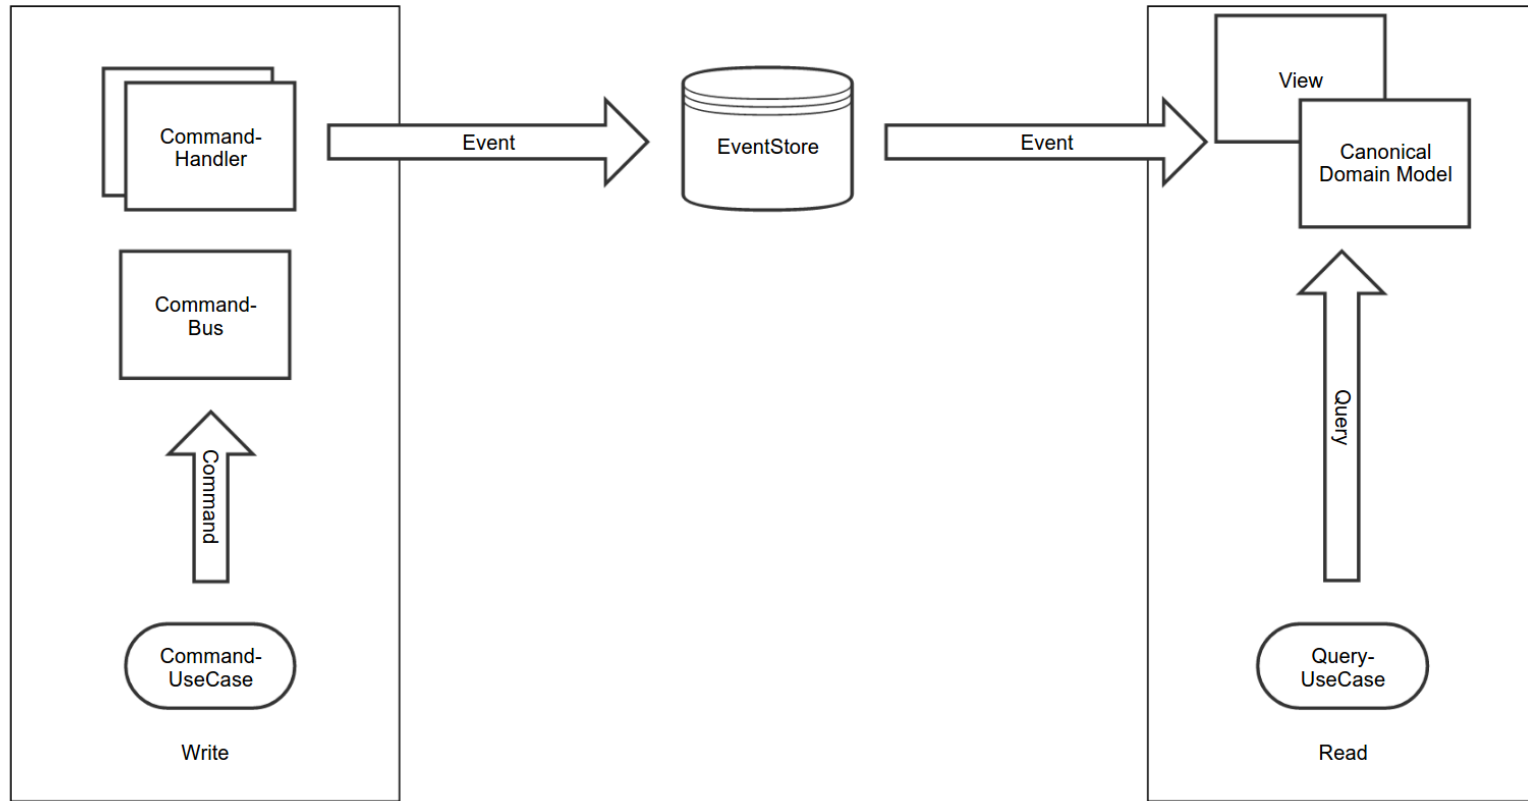
i want to **deposit** cash at the counter
in order to credit it to my account.

UseCase Withdrawal

UseCase Withdrawal

As a Customer

i want to **withdraw** money from my account at the counter
in order to cash it out.



What we learn

- Implement **write** side
 - minimal Commands
 - minimal CommandHandlers
 - minimal Events
- Implement **read** side
 - minimal EventHandlers (Views)
 - that populate the canonical Domain Model

```
git clone https://github.com/uweschaefer/es-basics.git
```

Session 1

1. Implement **ApplicationFacade.deposit/withdraw**
2. Create **Command classes** for both UseCases (see *CreateAccountCommand*)
3. Create **CommandHandlers** for both UseCases (see *CreateAccountHandler*)
4. Create **Event** classes for every UseCase (see *AccountCreatedEvent*)
5. Extend **AccountView** to aggregate Accounts
6. Pass the Tests

What just happened?


- complete Roundtrip
 - Command -> CommandHandler
 - Event -> EventHandler (View)
 - Query via DomainModel

EventSourcing Basics

Session 2

Dedicated QueryModel

Overview



Aggregation

AccountView = EventHandler that aggregates Events to an Account.

Not every Query within one Account

UseCase ValuedCustomerReport

UseCase ValuedCustomerReport

As a Manager, i want
a complete report the lists all *valued customers*
in order to free them from handling charges.

Specification ValuedCustomer

Specification ValuedCustomer

a valued Customer

has deposited an amount of **$\geq 1000\text{€}$ at least twice.**

How **not** to do that

Iterate Accounts and inspect their Depositions one by one.

What we learn

- Aggregate beyond Entity-Boundaries
- Create **dedicated Read Model**
 - find appropriate DataStructure
 - select Event-Types by @EventConsumer Methods

Session 2

1. `git clean -fd && git reset --hard`
session2
2. implement
ValuedCustomerReportView
3. pass the Tests

What just happened?

- we added a dedicated Read/Query-Model
 - beyond aggregate boundaries
 - Query-optimized Datastructure
 - PullViews have to actually **pull** the events from the ES at some point.

EventSourcing Basics
Session 3 (Bonuslevel)

Rolling Snapshot QueryModel

Overview

- ValuedCustomerSupport aggregates ALL depositions in the System
- gets slower & slower
- how to tackle that?

Problematic

every time a Report is needed, a new View has to be created.

```
facade.deposit(...);  
  
ValuedCustomerReportView report1 = new ValuedCustomerReportView(es);  
assertTrue(report1.isValuedCustomer(...));  
  
facade.deposit(...);  
  
ValuedCustomerReportView report2 = new ValuedCustomerReportView(es);  
assertTrue(report2.isValuedCustomer(...));
```

What if?

we could reuse a QueryModel, that is being updated, rather than re-created?

```
ValuedCustomerReportView report = new ValuedCustomerReportView(es);  
  
facade.deposit(...);  
assertTrue(report.isValuedCustomer(...));  
  
facade.deposit(...);  
assertTrue(report.isValuedCustomer(...));
```

discuss what would be necessary, conceptually?

No, really – **Discuss**

What we just learned?

- Concept of Rolling Snapshot

Session 3

1. git clean -fd && git reset --hard session3
2. look at *View.last*, *View.apply* and *PullView.pullEvents*
3. change ValueCustomerReport appropriately.
4. pass the Tests

What just happened?

- Rolling Snapshot
 - keeps latest
- Eventstore provides query of EventStream from **after** a particular event
- And yes, we can have more snapshots than one, if needed

Have a break.



Have a KitKat

EventSourcing Basics

Session 4

Event Design

UseCase Transfer

UseCase Transfer

As a user, i want to
transfer Money from my account to someone else's
in order to pay my rent online.

Acceptance Criteria

- *AccountUnknownException* if receiver or sender account does not exist
- *UnfundedTransferException* if sender does not have enough money (no debt allowed)

What we learn

- Event granularity matters
- Events need to reveal their intent
- Use of a ValidationModel
- Commands can be rejected

Session 4

1. `git clean -fd && git reset --hard session4`
2. Implement *TransferHandler*
3. pass the Tests

What just happened?

- Granularity: Events belong to ONE Aggregate
 - we need SendTransfer, RecieveTransfer
- Events reveal intent
 - do not reuse WithdrawnEvent etc, its a different UseCase!

EventSourcing Basics

Session 5 (Bonuslevel)

Dedicated WriteModel / ValidationModel

Command validation sometimes needs Context

Your take on Criteria 1?

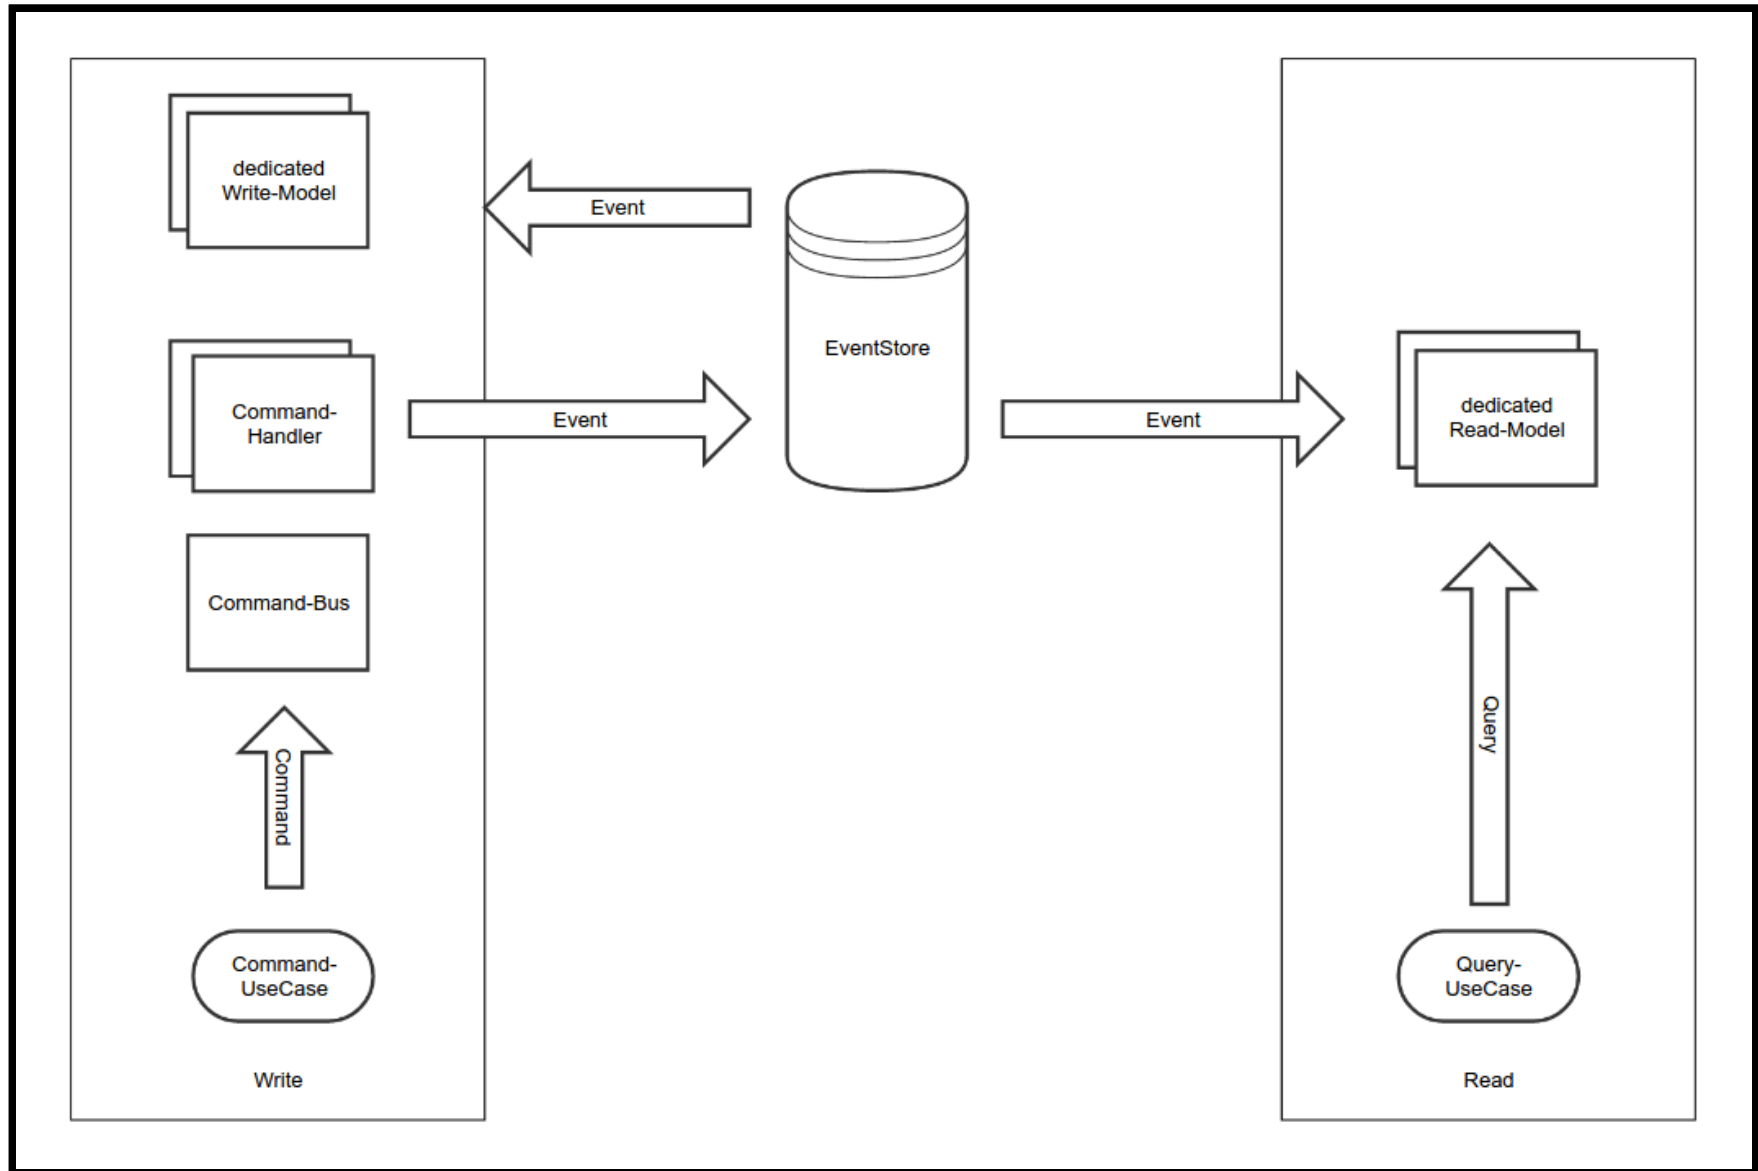
- *AccountUnknownException* if receiver or sender account does not exist

```
private boolean exists(UUID id) {  
    return repo.find(id) != null;  
}
```

Using AccountView just to find out, if an Account exists is wasteful.

All we really need to know
is if the aggregate exists.

Overview



CommandHandler

Responsibilities

- validate command
- accept or reject command based on that validation
- emit Messages on accepting

Session 5

1. `git clean -fd && git reset --hard session5`
2. implement `KnownAccountsView`
3. pass the Tests

What just happened?

- dedicated WriteModel / ValidationModel
- does not have to be DomainModel, as it does not need behaviour

EventSourcing Basics

Session 6

Side-Effects

- Some Commands may trigger external behavior.
- Replaying that would be problematic.

UseCase Notification

UseCase Notification

As a user, i want to
be notified by email when i recieve a transfer
in order to buy champagne asap.

What we learn

- how/where to model Side-Effect

Session 6

1. `git clean -fd && git reset --hard session6`
2. use `CreditNotificationService` to send mail
3. discuss where/how to do it properly
4. hint: see `CommandBus.publish()`
5. pass the Tests

What just happened?

- CommandBus has to be reliable
- Commands can be Effects, too
- Side-Effects can be modeled as Commands / CommandHandlers

Intermission

What about Consistency?

Did we relax Consistency compared to a normal CRUD/ORM implementation?

NO!

But where we could, how can we take advantage?

EventSourcing Basics
Session 7 (Bonuslevel)

Push-Views

Up to now, all views have been *PullViews*, that call *pullEvents()* to stream events into them.

Pro

- we can define when to update the View's State

Con

- we have to Query the EventStore in order to know, if View's State is stale
 - the more Queries we run, the more catastrophic this is:
 - bad Latency for Queries
 - high Contention on EventStore

UseCase GoldCustomers

UseCase GoldCustomers

As an accountant, i want to
know all the Gold-Customers
in order to be extra nice to them.

Specification Gold Customer

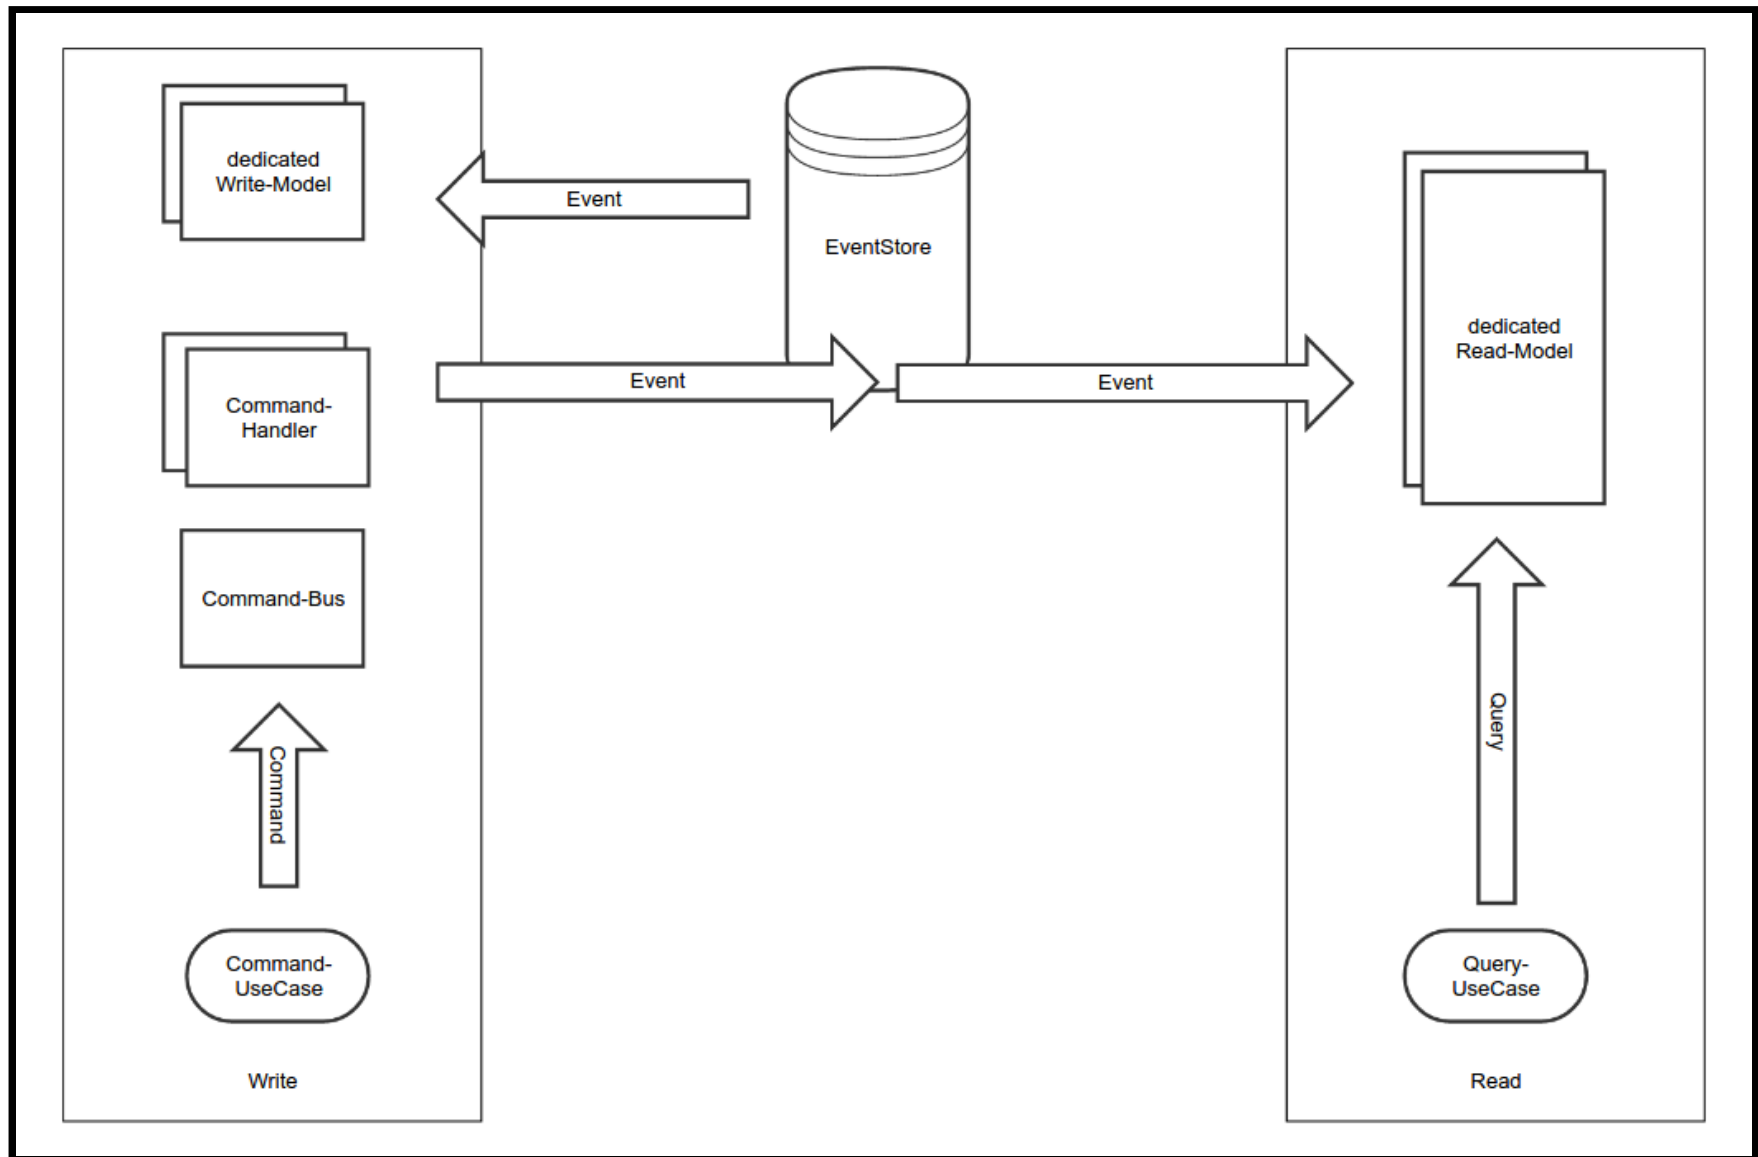
Specification Gold Customer

someone who recieved a transfer $\geq 10.000\text{€}$
at least once

Acceptance Criteria

- report must be **instant!** (low-latency)
- report must be a collection of Strings "<LASTNAME>, <FIRSTNAME>"
- order is not important
- only Transfers count – Depositions **must not** be examined
- report does not need to include GoldCustomer that recieved the status in the last few seconds...

... which means **Eventual Consistency** is ok



What would be necessary to push events to the view?

What we learn

- Use Push-Model for Views
- pros and cons of push vs pull

Session 7

1. git clean -fd && git reset --hard session7
2. implement GoldCustomers *extends*
PushView
3. pass the Tests

What just happened?

- Implemented a push-View that is updated by processing Events asynchronously
- Push reduces read latency
- introduces eventual consistency
- introduces concurrency
- PushViews mostly unusable as Validation Model (not strictly consistent)

One possible solution can be found here
`git clean -fd && git reset --hard theend`

Links

- O.Wolfs CQRS Slides
 - <https://speakerdeck.com/owolf/cqrs-for-great-good-2>
- Greg Young's Blog
 - <https://goodenoughsoftware.net/>
- Axon mature ES Framework
 - <http://www.axonframework.org/>
- Lagom Modern ES Framework based on Akka
 - <https://www.lightbend.com/lagom>
- Microsofts CQRS/ES Patterns & Practices
 - <https://msdn.microsoft.com/en-us/library/jj554200.aspx>

Q & A